

Microprocessor Design Trends

- Joy's Law [Bill Joy of BSD4.x and Sun fame]

$$\text{MIPS} = 2^{\text{year}-1984}$$

- Millions of instructions per second [MIPS] executed by a single chip microprocessor
- More realistic rate is a doubling of MIPS every 18 months [or a quadrupling every 3 years]
- What ideas and techniques in new microprocessor designs have contributed to this continued rate of improvement?

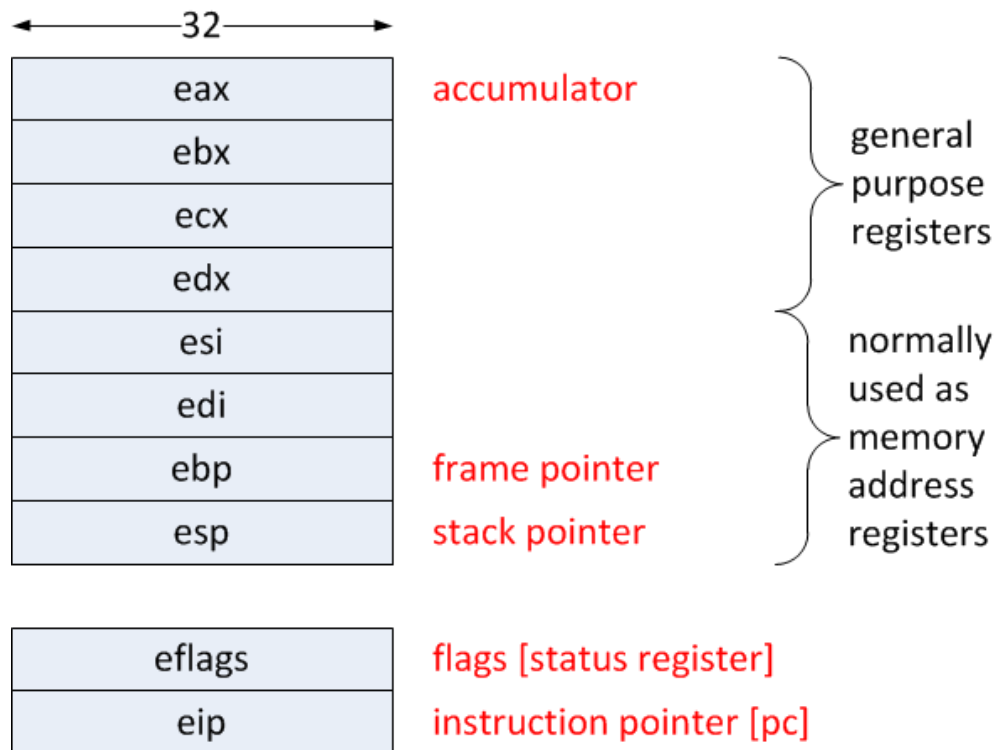
Some of the ideas and techniques used...

- smaller VLSI feature sizes [1 micron (μ) -> 10nm]
- increased clock rate [1MHz -> 4GHz]
- reduced vs complex instruction sets [RISC vs CISC]
- Faster memory access modes (eg burst)
- integrated on-chip MMUs, FPU, ...
- pipelining
- superscalar [multiple instructions/clock cycle]
- multi-level on-chip instruction and data caches
- streaming SIMD [single instruction multiple data] instruction extensions [MMX, SSE]
- multiprocessor support
- hyper threading and multi core
- direct programming of graphics co-processor
- high speed point to point interconnect [Intel QuickPath, AMD HyperTransport]
- solid state disks
- ...

IA32 [Intel Architecture 32 bit]

- IA32 first released in 1985 with the 80386 microprocessor
- IA32 still used today by current Intel CPUs
- modern Intel CPUs have many additions to the original IA32 including MMX, SSE1, SSE2, SSE3, SSE4 and SSE5 [Streaming SIMD Extensions] and even an extended 64 bit instruction set when operating in 64 bit mode [named IA-32e or IA-32e or x64]
- 32 bit CPU [performs 8, 16 and 32 bit arithmetic]
- 32 bit virtual and physical address space 2^{32} bytes [4GB]
- each instruction a multiple of bytes in length [1 to 17+]


Registers [far fewer than a typical RISC]



NB: floating point and SSE registers, ... not shown

Registers...

- "e" in eax = extended = 32bits



- possible to access 8 and 16 bit parts of eax, ebx, ecx and edx using alternate register names

Instruction Format

- two address [**will use Microsoft assembly language syntax used by VC++, MASM**]

add ←
 eax, ebx ; eax = eax + ebx [**right to left**]

- alternative gnu syntax*

addl →
 %ebx, %eax ; eax = eax + ebx [**left to right**]

- two operands normally

←
register/register
register/immediate
register/memory
memory/register

- memory/memory and memory/immediate are NOT allowed

Supported Addressing Modes

$[a]$ = contents of
memory address a

<i>addressing mode</i>	<i>example</i>	
immediate	<code>mov eax, n</code>	$eax = n$
register	<code>mov eax, ebx</code>	$eax = ebx$
direct/absolute	<code>mov eax, [a]</code>	$eax = [a]$
indexed	<code>mov eax, [ebx]</code>	$eax = [ebx]$
indexed	<code>mov eax, [ebx+n]</code>	$eax = [ebx + n]$
scaled indexed	<code>mov eax, [ebx*s+n]</code>	$eax = [ebx*s + n]$
scaled indexed	<code>mov eax, [ebx+ecx]</code>	$eax = [ebx + ecx]$
scaled indexed	<code>mov eax, [ebx+ecx*s+n]</code>	$eax = [ebx + ecx*s + n]$

- address computed as the sum of a register, a scaled register and a 1, 2 or 4 byte signed constant n ; can use most registers
- scaled indexed addressing used to index into arrays
- scaling constant s can be 1, 2, 4 or 8

Assembly Language Tips

- size of operation can often be determined implicitly by MASM, but when unable to do so, size needs to be specified explicitly

mov eax, [ebp+8] ; implicitly 32 bit [as eax is 32 bits]

mov ah, [ebp+8] ; implicitly 8 bit [as ah is 8 bits]

dec [ebp+8] ; decrement memory location [ebp+8] by 1
; assembler unable to determine operand size
; is it an 8, 16 or 32 bit value??

dec DWORD PTR [ebp+8] ; make explicitly 32 bit

dec WORD PTR [ebp+8] ; make explicitly 16 bit

dec BYTE PTR [ebp+8] ; make explicitly 8 bit

NB: unusual assembly language syntax

Assembly Language Tips...

- memory/immediate operations NOT allowed

~~mov [ebp+8], 123~~ ; NOT allowed and operation size ALSO unknown

mov eax, 123 ; use 2 instructions instead...
mov [ebp+8], eax ; implicitly 32 bits

- lea [**load effective address**] is a useful instruction for performing simple arithmetic

lea eax, [ebx+ecx*4+16] ; $eax = ebx + ecx * 4 + 16$

Basic Instruction Set

mov	move
xchg	exchange
add	add
sub	subtract
cdq	convert double to quadword
imul	signed multiply
mul	unsigned multiply
inc	increment by 1
dec	decrement by 1
neg	negate
cmp	compare
lea	load effective address
test	AND operands and set flags
and	and
or	or
xor	exclusive or
not	not

push	push onto stack
pop	pop from stack
sar	shift arithmetic right
shl	shift logical left
shr	shift logical right
jmp	unconditional jump
j {e, ne, l, le, g, ge}	signed jump
j {b, be, a, ae}	unsigned jump
call	call subroutine
ret	return from subroutine

- should be enough instructions to complete tutorials
- Google [Intel® 64 and IA-32 Architectures Software Developer's Manual 2A, 2B, 2C](#) for details

Assembly Language Tips...

- quickest way to clear a register?

```
xor    eax, eax
```

; exclusive OR with itself

```
mov    eax, 0
```

; instruction occupies more bytes and...
; probably takes longer to execute

- quickest way to test if a register is zero?
- NB mov instruction doesn't update condition code flags

```
test   eax, eax
```

```
je     ...
```

; AND eax with itself, set flags and...
; jump if zero

Function Calling

reminder of the steps normally carried out during a function/procedure call and return

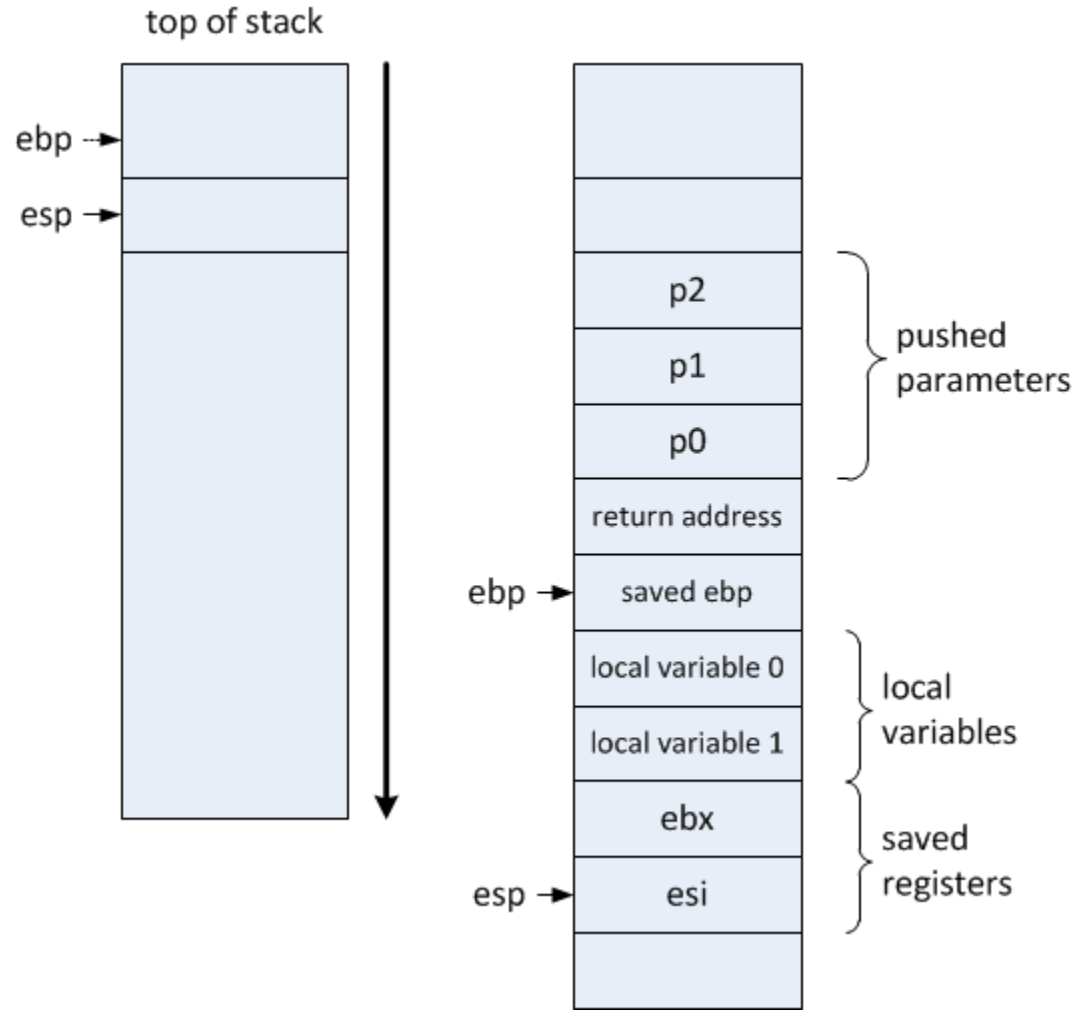
- pass parameters [**evaluate and push on stack**]
- enter new function [**push return address and jump to first instruction of function**]
- allocate space for local variables [**on stack by decrementing esp**]
- save registers [**on stack**]

<function body>

- restore saved registers [**from stack**]
- de-allocate space for local variables [**increment esp**]
- return to calling function [**pop return address from stack**]
- remove parameters [**increment esp**]

IA32 Function Stack Frame

- stack frame after call to `f(p0, p1, p2)`
- stack grows down in memory [from highest address to lowest]
- parameters pushed right to left
- NB: stack always aligned on a 4 byte boundary [it's not possible to push a single byte]
- `ebp` used as a frame pointer
parameters and locals accessed relative to `ebp` [eg `p0 @ ebp+8`]



IA32 Calling Conventions

- several IA32 procedure/function calling conventions
- use Microsoft `_cdecl` calling convention [as per previous diagram] so C/C++ and IA32 assembly language code can mixed

function result returned in `eax`

`eax`, `ecx` and `edx` considered volatile and are NOT preserved across function calls

caller removes parameters

- why are parameters pushed right-to-left??

C/C++ pushes parameters right-to-left so functions like `printf(char *formats, ...)` [which can accept an arbitrary numbers of parameters] can be handled more easily since the first parameter is always stored at `[ebp+8]` irrespective of how many parameters are pushed

Accessing Parameters and Local Variables

- ebp used as a frame pointer; parameters and local variables accessed at offsets from ebp
- can avoid using a frame pointer [**normally for speed**] by accessing parameters and locals variables relative to the stack pointer, but more difficult because the stack pointer can change during execution [**BUT easy for a compiler to track**]
- parameters accessed with +ve offsets from ebp [**see stack frame diagram**]

p0 @ [ebp+8]

p1 @ [ebp+12]

...

- local variables accessed with -ve offsets from ebp [**see stack frame diagram**]

local variable 0 @ [ebp-4]

local variable 1 @ [ebp-8]

...

Consider the IA32 Code for a Simple Function

```
int f (int p0, int p1, int p2) {    // parameters
    int x, y;                      // local variables
    x = p0 + p1;
    ...
    return x + y;                  // result
}
```

- a call `f(p0, p1, p2)` matches stack frame diagram on previous slide
- 3 parameters *p0*, *p1* and *p2* and 2 local variables *x* and *y*
- need to generate code for
 - *calling function f*
 - *function f entry*
 - *function f body*
 - *function f exit*

IA32 Code to Call Function f

- parameters $p0$, $p1$ and $p2$ pushed onto stack by caller right to left

f(1, 2, 3)

```
push    3
push    2
push    1
call    f
add     esp, 12
```

; push immediate values...


; right to left

;

; call f

; add 12 to esp to remove parameters from stack

push return address
and jump to f



Function Entry

- need instructions to save ebp [**old frame pointer**] and ...
- initialize ebp [**new frame pointer**] and ...
- allocate space for local variables on stack and ...
- push non volatile registers used by function onto stack

```
f:    push    ebp                ; save ebp
      mov     ebp, esp          ; ebp -> new stack frame
      sub     esp, 8            ; allocate space for locals x and y
      push    ebx                ; save non volatile registers used by function

      <function body>          ; function body

      <function exit>         ; function exit
```

NB: cdecl convention means there is NO need to save eax, ecx and edx

Function Body

- parameters pushed on stack and ...
- space already allocated for local variables

parameters p0 @ [ebp+8] and p1 @ [ebp+12]

locals x @ [ebp-4] and y @ [ebp-8]

- $x = p0 + p1$

```
mov    eax, [ebp+8]    ; eax = p0
add    eax, [ebp+12]   ; eax = p0 + p1
mov    [ebp-4], eax    ; x = p0 + p1
```

- return $x + y$;

```
mov    eax, [ebp-4]    ; eax = x
add    eax, [ebp-8]    ; eax = x + y
```

NB: result returned in eax

Function Exit

- need instructions to unwind stack frame at function exit

```
...  
pop      ebx           ; restore saved registers  
mov       esp, ebp      ; restore esp  
pop       ebp           ; restore previous ebp  
ret       0             ; return from function
```

- ret pops return address from stack and...
- adds integer parameter to esp [*used to remove parameters from stack*]
- if integer parameter not specified, defaults to 0
- since using *_cdecl* convention caller will remove parameters from stack
- make sure you know why a stack frame needs to be created for each function call

IA32 Code for Accessing an Array

```
int a[100];                // global array of int
```

```
main(...) {  
    a[1] = a[2] + 3;        // constant indices  
}
```

- int is 4 bytes
- assume array a is stored at absolute address a
- a[0] store at address a, a[1] at a+4, a[2] at a+8, a[n] at a+n*4

```
mov    eax, [a+8];         // eax = a[2]  
add    eax, 3              // eax = a[2] + 3  
mov    [a+4], eax          // a[1] = a[2] + 3
```

IA32 Code for Accessing an Array ...

```
int *a = (int*) malloc(100*sizeof(int));    // THIS TIME array allocated on heap
```

```
int p() {  
    int i = ...;           // local variable i @ [ebp-4]  
    int j = ...;           // local variable j @[ebp-8]  
    ...  
    a[i] = a[j] + 3;        // variable indices  
}
```

- global variable a contains the address of the array allocated on heap

```
mov    edx, [a]            // edx -> a  
mov    eax, [ebp-8]        // eax = j  
mov    eax, [edx+eax*4]    // eax = a[j]  
add    eax, 3              // eax = a[j]+3  
mov    ecx, [ebp-4]        // ecx = i  
mov    [edx+ecx*4], eax    // a[i] = a[j]+3
```

Putting it Together – Tutorial 1...

- Mixing C/C++ and IA32 Assembly Language
- Example using Visual Studio 2010/2013/2015/2017, VC++ and MASM
- t1Test.cpp calls an assembly language versions of fib [as a demonstration], min, p and gcd
- create a VC++ Win32 Console Application [call it t1Test]
- select project name (t1Test), click on Project menu, select "Build Customizations..." and tick masm
- add files fib32.h, fib32.asm, t1.h and t1.asm
- Right click on .asm files make sure [Properties][General][Item Type] is set to Microsoft Macro Assembler
- You can create the t.h and t1.asm externally and include them into the project using [Project][Add Existing Item...]

Putting it Together – Tutorial 1

- make sure the configuration is x86 [you can delete the default x64 configuration as it is not applicable]
- how to see the code generated by the VC++ compiler??
 - right click on C/C++ file name [Properties] [C/C++] [Output Files] [Assembler Output] and select Assembly, Machine Code and Source [listing has a .cod extension]
 - code generated in Debug and Release mode will be different
- setting breakpoints in .asm file
 - assembly source code debugging hasn't worked properly since VS2013
 - to debug min (for example), set breakpoint in .cpp file before call to min
 - when breakpoint reached, select [Debug][Windows][Disassembly]
 - THEN single step using F11
 - hover mouse over register names to see their values etc.

Putting it Together – Tutorial 1...

fib32.h

- declare fib_IA32a(int) and fib_IA32b(int) as external C functions so they can be called from a C/C++ program

```
extern "C" int g;                // external global int
extern "C" int _cdecl fib_IA32a(int); // external function
```

- extern "C" because C++ function names have extra characters which encode their result and parameter types

fib32.asm

- fib_IA32a(int) – *mechanical* code generation simulating Debug mode
- fib_IA32b(int) – *optimized* code generation simulating Release mode
- MASM specific directives at start of file
- .data and .code sections
- public

Putting it Together – Tutorial 1...

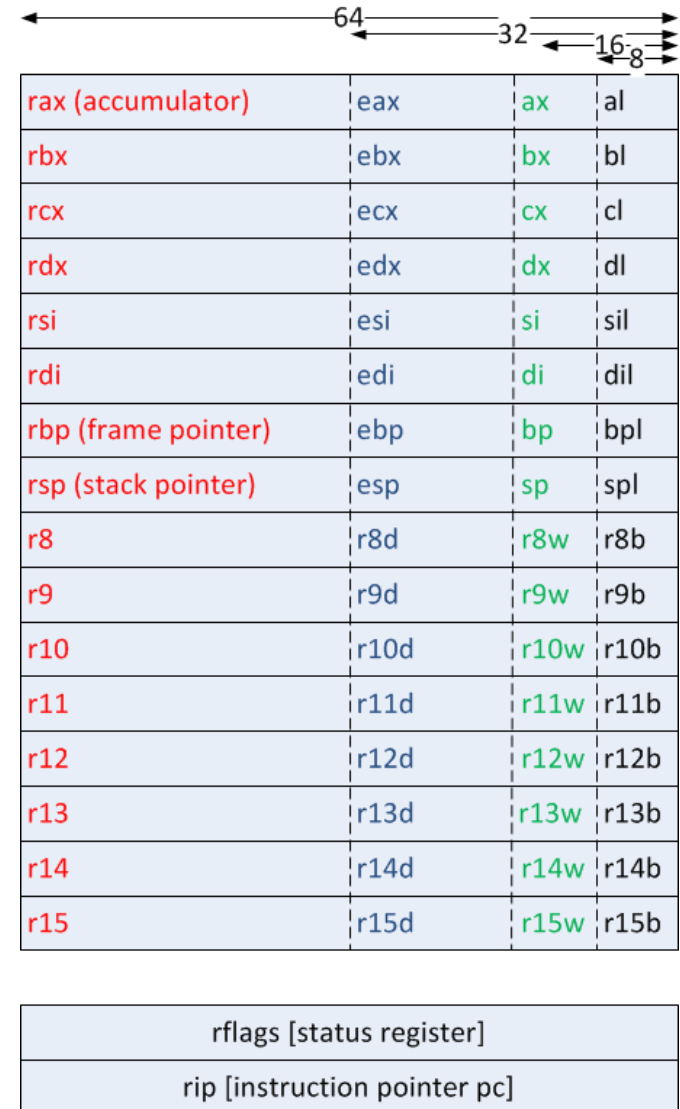
- t1Test.cpp [**_tmain**]
- #include fib32.h and t1.h
- calls fib_IA32a(n) and fib_IA32b(n) like any other C/C++ function
- file also contains
 - 1) a C++ version of fib(n) and...
 - 2) a version of fib(n) that mixes C/C++ and IA32 assembly language using the IA32 inline assembler supported by the VC++ compiler
- call ALL versions of fib(n) for n = 1 to 20
- Visual Studio automatically compiles t1Test.cpp, assembles fib32.asm and t1.cpp and links them to produce an executable which can be run
- WARNING: Visual Studio on SCSS machines (eg ICT Huts) has problems when source files are stored on a Network drive

IA32 AND X64

IA32 AND x64

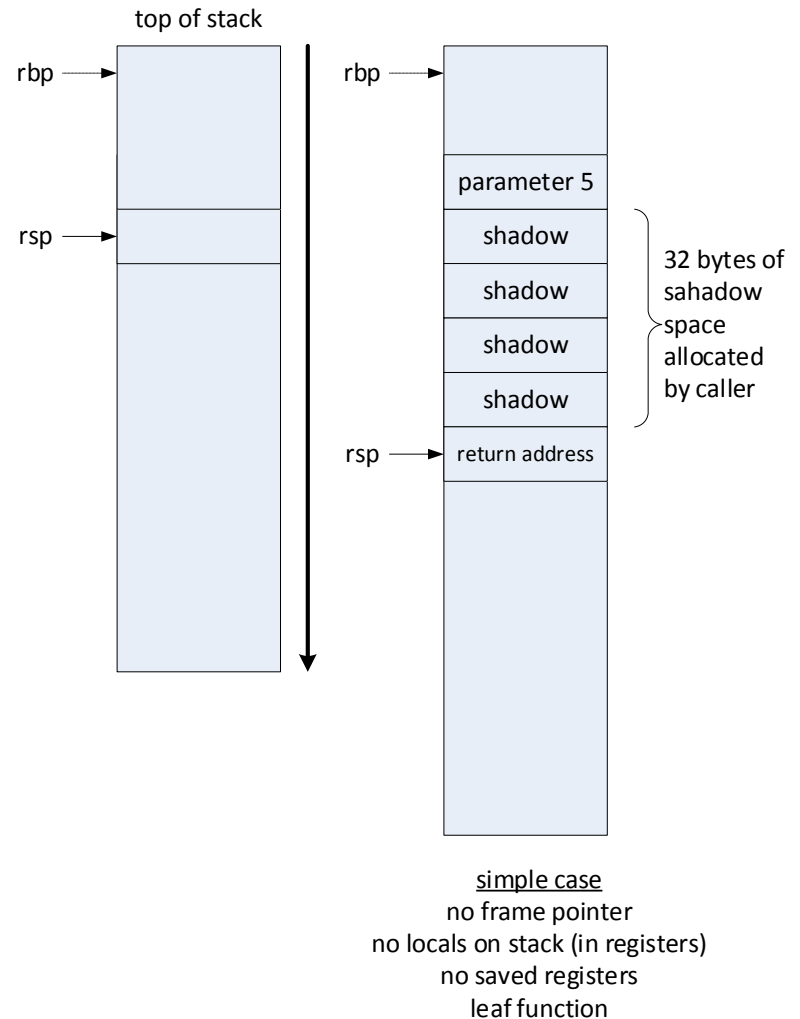
x64 Basics

- extension of IA32
- originally developed by AMD
- IA32 registers extended to 64 bits rax ... rsp, rflags and rip
- 8 additional registers r8 .. r15
- 64, 32, 16 and 8 bit arithmetic
- *same* instruction set
- 64 bit virtual and physical address spaces
[theoretically anyway]
- $2^{64} = 16 \text{ Exabytes} = 16 \times 10^{18} \text{ bytes}$



x64 Function Calling

- use Microsoft calling convention
- first 4 parameters passed in rcx, rdx, r8 and r9 respectively
- additional parameters passed on stack [**right to left**]
- stack always aligned on an 8 byte boundary
- caller must allocate 32 bytes of *shadow space* on stack
- rax, rcx, rdx, r8, r9, r10 and r11 volatile
- having so many registers often means:
 1. can use registers for local variables
 2. no need to use a frame pointer
 3. no need to save/restore registers



x64 Function Calling Microsoft Convention ...

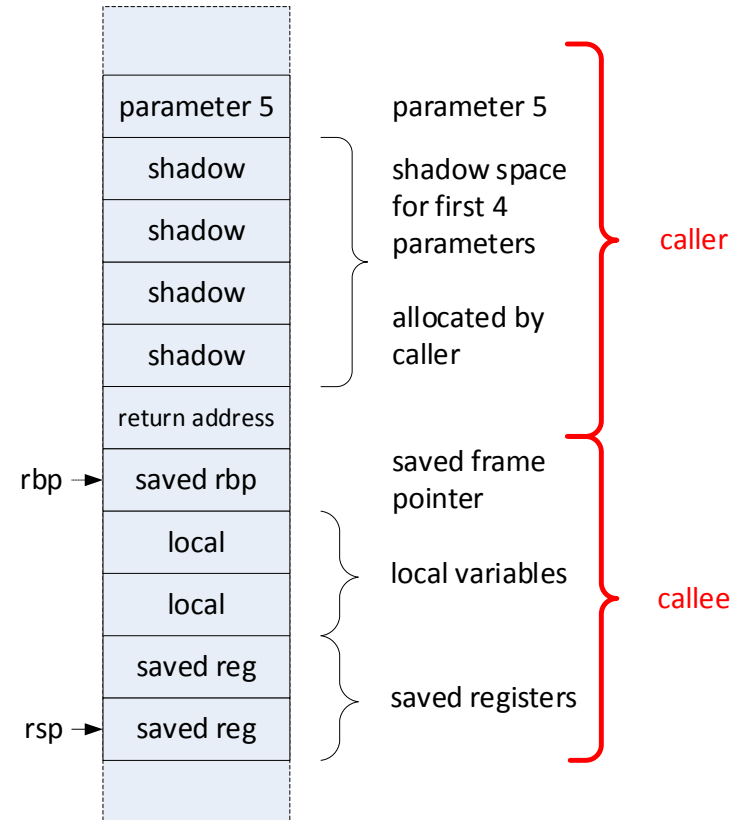
- caller must allocate 32 bytes (4 x 8bytes) of *shadow space* on the stack before calling a function [**regardless of the actual number of parameters used**] and to deallocate the *shadow space* afterwards
- called functions can use its *shadow space* to spill rcx, rdx, r8, and r9 [**spill = save in memory**]
- called functions may use the *shadow space* for any purpose whatsoever and consequently may read and write to it as it sees fit [**which is why it needs to be allocated**]
- 32 bytes of *shadow space* must be made available to all functions, even those with fewer than four parameters
- what are the advantages of having shadow space?

x64 Function Calling Unix/Linux

- brief description
- first six parameters passed in registers RDI, RSI, RDX, RCX, R8, R9 respectively
- additional arguments are passed on the stack [right to left]
- use of frame pointer [rbp], allocation of locals on stack and saving of registers as per Microsoft convention
- result returned in rax
- registers ebp, rbx, r12, r13, r14 and r15 non volatile
- no shadow space as per Microsoft convention

x64 Function Calling (Microsoft Convention) ...

- a more complex x64 stack frame
- callee has 5 parameters, so parameter 5 passed on stack
- parameters 1 to 4 passed in rcx, rdx, r8 and r9
- shadow space allocated
- old frame pointer saved and new frame pointer initialised [**rbp**]
- space allocated for local variables on stack [**if needed**]
- registers saved on stack [**if needed**]



x64 Function Calling (Microsoft Convention) ...

```
_int64 fib(_int64 n) {  
    INT64 fi, fj, t;
```

```
    if (n <= 1)  
        return n;
```

```
    fi = 0; fj = 1;  
    while (n > 1) {  
        t = fj;  
        fj = fi + fj;  
        fi = t;  
        n--;
```

```
    }  
    return fj;  
}
```

- use `_int64` to declare 64 bit integers [**Microsoft specific**]
- alternatively

declare 64 bit integers using `long long`

```
#define INT64 long long
```

- parameter `n` passed to function in `rcx`
- leaf function [**as fib doesn't call any other functions**]
- usually easier to code with x64 assembly language rather than IA32 because a simpler stack frame is used and more registers are available

x64 Function Calling (Microsoft Convention) ...

```
fib_x64:    mov     rax, rcx           ; rax = n
            cmp     rax, 1             ; if (n <= 1)
            jle     fib_x64_1         ; return n
            xor     rdx, rdx           ; fi = 0
            mov     rax, 1             ; fj = 1
fib_x64_0:  cmp     rcx, 1             ; while (n > 1)
            jle     fib_x64_1         ;
            mov     r10, rax          ; t = fj
            add     rax, rdx           ; fj = fi + fj
            mov     rdx, r10          ; fi = t
            dec     rcx               ; n--
            jmp     fib_x64_0         ;
fib_x64_1:  ret                       ; return
```

- code ONLY uses volatile registers
- leaf function so no need to allocate shadow space

x64 Function Calling (Microsoft Convention) ...

```
_int64 xp2(_int64 a, _int64 b) {  
    printf("a = %l64d b = %l64d a+b = %l64d\n", a, b, a + b);  
    return a + b;    // NB  
}
```

- uses %l64d to format a 64 bit integer
- parameters **a** and **b** passed to xp2 in rcx and rdx respectively
- need to call external printf(...) function with 4 parameters

```
rcx [address of format string]  
rdx [a]  
r8 [b]  
r9 [a+b]
```

x64 Function Calling (Microsoft Convention) ...

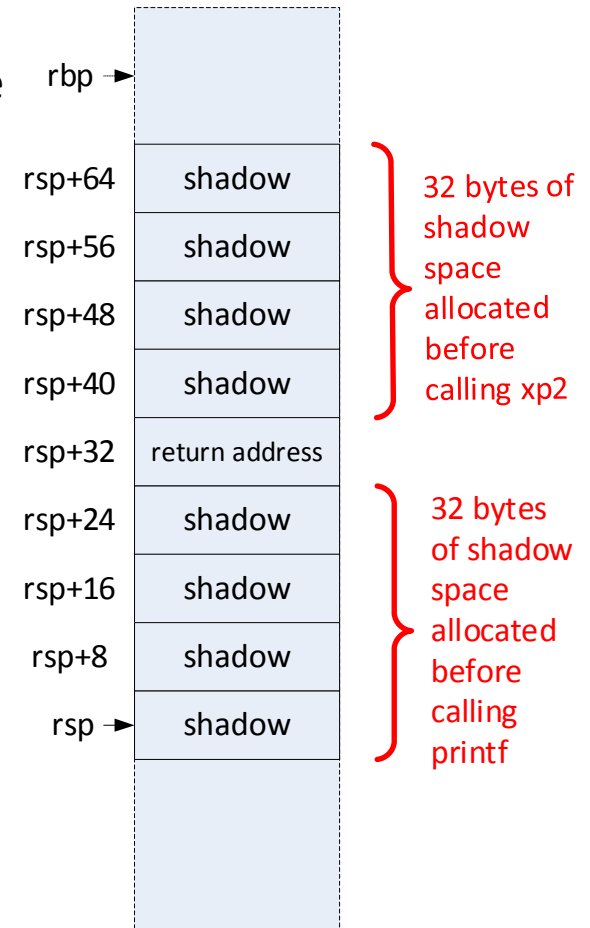
fxp2 db 'a = %l64d b = %l64d a+b = %l64d', 0AH, 00H ; ASCII format string

```
xp2: push    rbx                ; save rbx
      sub     rsp, 32           ; allocate shadow space
      lea     r9, [rcx+rdx]     ; printf parameter 4 in r9 {a+b}
      mov     r8, rdx           ; printf parameter 3 in r8 {b}
      mov     rdx, rcx          ; printf parameter 2 in rdx {a}
      lea     rcx, fxp2         ; printf parameter 1 in rcx {&fxp2}
      mov     rbx, r9           ; save r9 in rbx so preserved across call to printf
      call    printf           ; call printf
      mov     rax, rbx          ; function result in rax = rbx {a+b}
      add     rsp, 32           ; deallocate shadow space
      pop     rbx              ; restore rbx
      ret                     ; return
```

x64 Function Calling (Microsoft Convention) ...

- instead of using rbx to preserve r9 across the call to printf, an alternate approach is to use a location its shadow space [eg. `rsp+64`]

```
xp2: sub    rsp, 32      ; allocate shadow space
      lea    r9, [rcx+rdx] ; printf parameter 4 in r9 {a+b}
      mov    r8, rdx      ; printf parameter 3 in r8 {b}
      mov    rdx, rcx     ; printf parameter 2 in rdx {a}
      lea    rcx, fxp2    ; printf parameter 1 in rcx
      mov    [rsp+64], r9 ; save r9 in shadow space so...
      call   printf       ; preserved across call to printf
      mov    rax, [rsp+64] ; result in rax = saved r9 {a+b}
      add    rsp, 32      ; deallocate shadow space
      ret                ; return
```



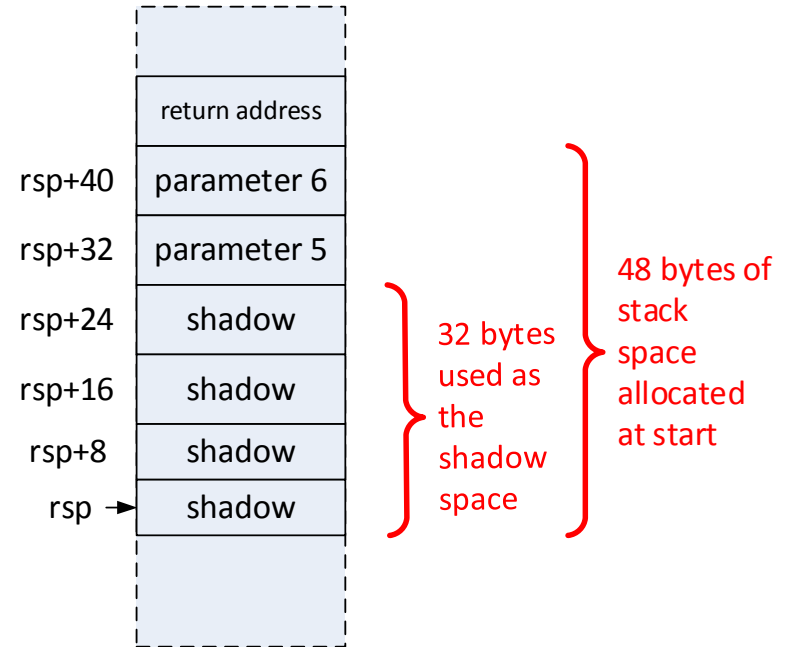
x64 Function Calling (Microsoft Convention) ...

Typical code generation strategy

- shadow space allocated ONCE at start of function
- allocate enough stack space to accommodate calls to the function with the most parameters [NB: must allocate a minimum 32 bytes for the shadow space]
- use the same stack space [and registers] to pass parameters to ALL the functions it calls
- straightforward for compiler to determine how much stack space is required

Typical code generation strategy...

```
function f(...)
...
printf(5 parameters);
...
printf(6 parameters);
...
printf(2 parameters);
..
}
```

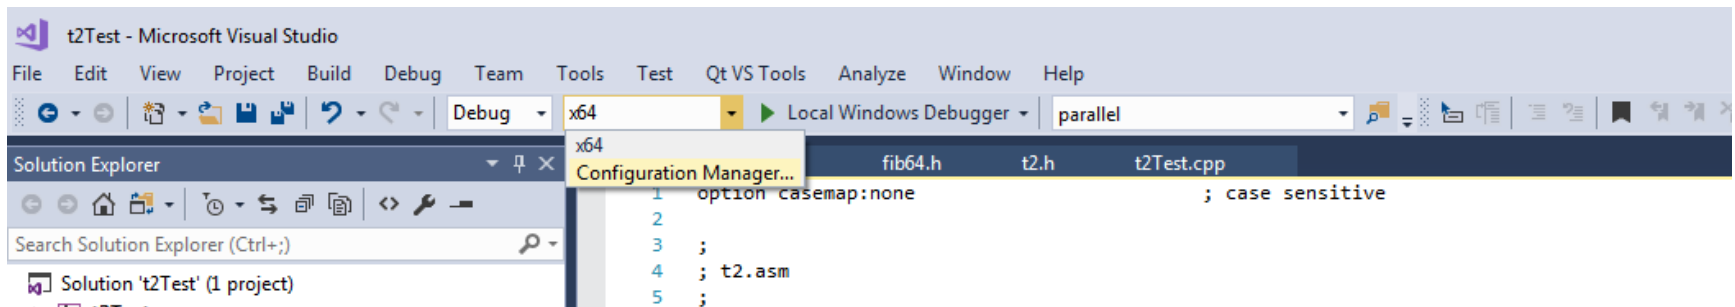


- maximum number of parameters is 6
- need to allocate $6 \times 8 = 48$ bytes on stack at start
- in general, allocate $\max(32, n \times 8)$ bytes where n is the maximum number of parameters (minimum 32 bytes of stack space for shadow space allocated)
- parameter 5 is moved directly to stack (NOT pushed) eg `mov [rsp+32], eax`
- reuse allocated stack space for all 3 calls to `printf`
- deallocate stack space on exit

IA32 AND x64

Using Visual Studio

- fib64.h, fib64.asm and t2Test.cpp available on CS3021/CS3421 website
- need to create a console application and use the Configuration Manager to select a x64 solution platform



- one way to link with printf is to include the following at the head of t2.asm

```
includelib legacy_stdio_definitions.lib
extrn      printf:near
.data
```

- no x64 inline assembler, can use intrinsics defined in [intrin.h](#) instead

Summary

- you are now able to:
 - write simple IA32 assembly language functions
 - write simple x64 assembly language functions
 - call IA32/x64 assembly language functions from C/C++
 - program the two most widely used CPUs in assembly language
[IA32/x64 and ARM]

