

Concurrency with Pi

Lecture 27

CS 565

(slides adapted from Martin Abadi)

1

Concurrency



Concurrency comes in several flavors:

- Distributed processes
 - Code running at multiple sites
- Threads
 - A language mechanism for specifying interleaving computations;
often run on a single processor
- Parallel (SIMD)
 - A single program but with simultaneous operations on multiple
data

Different notions) different research communities

- Internet agents vs. high-performance physics calculations

Ways to Describe Concurrency



Sequential processes are modeled by λ -calculus

- inevitable: the most natural way to observe an algorithm is to examine its output for various inputs) functions

A concurrent system is naturally non-deterministic

- Interleaving of atomic actions from different processes

Concurrent processes can be observed in many ways

- When are two concurrent systems equivalent ?
- Intra-process behavior vs. inter-process behavior

Concurrency can be described in many ways

- Process creation: fork/wait, cobegin/coend, data parallelism
- Process communication: shared memory, message passing
- Process synchronization: monitors, semaphores, transactions

This lead to a variety of process calculi

Communication Primitives



As usual we focus on foundations.

- Focus on communication as fundamental concept.
- Not a foundation for everything.

Not much to say about parallel computing here.


Communication through shared variables.

- Found mostly in parallel computing and threads.

Communication through messages.

- synchronous or asynchronous
- static or dynamic communication topology
- first-order or high-order data
- local or distributed

Communication in Languages



Historically the treatment of communication is weak.

- I/O often not considered part of the language.
- Added only as an afterthought (Pascal)

Even “modern” languages have primitive language I/O facilities.

- Messages are rare.
- Higher-level remote procedure/method call is rare.

Focus on communication by message passing.

Calculi and Languages




Several calculi and languages rely on message-passing:

- Communicating Sequential Processes (CSP) (Hoare, 1978)
- Occam (Jones)
- Calculus of Communicating Systems (CCS) (Milner, 1980)
- The Pi calculus (Milner, 1989 and others)
- Pict (Pierce and Turner)
- Concurrent ML (Reppy)
- Java RMI

Sometimes messaging is built in higher-level primitives

- Remote procedure call
- Remote method invocation

The Pi Calculus



The pi calculus is a process algebra (a la CCS)

Constructs for concurrency.

Communication on channels

- channels are first-class
 - channel names can be sent on channels
- access restrictions for channels

In λ -calculus everything is a function

In Pi calculus everything is a process

Communication in the Pi Calculus



Processes communicate on channels:

$c\langle M \rangle$ send message M on channel c .

$c(x)$ receives x on channel c .

Sequencing:

$c\langle M \rangle . p$ sends message M on c , then does p .

$c(x) . p$ receives x on c , then does p with x .

Concurrency:

$p \mid q$ is the parallel composition of p and q .

Replication:

$! p$ creates an infinite number of replicas of p

Examples



For example we might define

```
Speaker    = air<M>
Phone      = air(x).wire<x>
ATT        = wire(x).fiber<x>
System     = Speaker | Phone | ATT
```

Communication between processes is modeled by reduction:

```
Speaker | Phone → wire<M>
wire<M> | ATT  → fiber<M>
```

Composing these reductions we get:

```
Speaker | Phone | ATT → fiber<M>
```

Channel Visibility



Anybody can monitor an unrestricted channel:

Consider that we define

```
WireTap = wire(x).wire<x>.NSA<x>
```

- Copies the messages from the wire to NSA
- Possible since the name “wire” is globally visible

Now

```
WireTap | wire<M> | ATT !
wire<M>.NSA<M> | ATT !
NSA<M> | fiber<M>                                OOPS !
```

Restriction



The restriction operator $(\nu c)p$ makes a fresh channel c within process p .

- ν is the Greek letter “nu”
- The name “ c ” is local (bound) in p

Restricted channels cannot be monitored.

```
wire(x) ... | ( $\nu$  wire)(wire<M> | ATT) →  
wire(x) ... | fiber<M>
```

The scope of the name “ $wire$ ” is restricted

There is no conflict with the global $wire$

Restriction and Scope



Restriction

- is a binding construct (like λ)
- is lexically scoped
- allocates a new object (a channel)

$(\nu c)p$ is like `let c = new Channel() in p`

In particular, c can be sent outside its scope.

- But only if p decides so

First-Class Channels



A channel c can leave its scope of declaration

- via a message $d\langle c \rangle$ from within p

Allowing channels to be sent as messages means communication topology is dynamic.

- If channels are not sent as messages (or stored in the heap) then the communication topology is static.
- This differentiates Pi-calculus from CCS

Example of First-Class Channels



Consider:

```
MobilePhone = air(x).cell<x>
ATT1         = wire<cell>
ATT2         = wire(y).y(x).fiber(x)
```

in

```
(v cell)( MobilePhone | ATT1 ) | ATT2
```

ATT1 is trying to pass c out of the static scope of the restriction
 $v \text{ cell}$

Scope Extrusion



A channel is a name.

- First-class names must be usable even outside their original scope.

The pi calculus allows restrictions to move:

$$((\nu c)p) \mid q = (\nu c)(p \mid q) \quad \text{if } c \text{ not free in } q$$

Renaming is needed in general:

$$\begin{aligned} ((\nu c)p) \mid q &= \\ ((\nu d)[d/c]p) \mid q &= \\ (\nu d)([d/c]p \mid q) &\quad \text{where } d \text{ is fresh (not in } p \text{ or } q) \end{aligned}$$

Example, Continued



$$\begin{aligned} (\nu \text{ cell})(\text{MobilePhone} \mid \text{ATT1}) \mid \text{ATT2} &= \\ (\nu \text{ cell})(\text{MobilePhone} \mid \text{ATT1} \mid \text{ATT2}) &\rightarrow \\ (\nu \text{ cell})(\text{MobilePhone} \mid \text{cell}(x).\text{fiber}\langle x \rangle) \end{aligned}$$

Scope extrusion distinguishes the pi calculus from other process calculi.

Syntax of the Pi Calculus



There are many versions of the Pi calculus

A basic version:

```
p, q ::=
0      nil process
x<y>.p  sending
x(y).p  receiving
p | q   parallel composition
!p      replication
(ν x)p  restriction
```

Note that only variables can be channels and messages

Operational Semantics



One basic rule of computation

$$\frac{}{x\langle y \rangle.p \mid x(z).q \rightarrow p \mid [y/z]q}$$

- Synchronous communication between a sender and a receiver
- Both the sender and the receiver proceed afterwards

Rules for identifying senders and receivers

$$\frac{p \rightarrow p'}{p \mid q \rightarrow p' \mid q} \qquad \frac{p \rightarrow p'}{(\nu x)p \rightarrow (\nu x)p'}$$

$$\frac{p \equiv p' \quad p' \rightarrow q' \quad q' \equiv q}{p \rightarrow q}$$

Structural Congruence



$$\frac{}{p \equiv p} \quad \frac{q \equiv p}{p \equiv q} \quad \frac{p \equiv q \quad q \equiv r}{p \equiv r}$$

$$\frac{p \equiv p'}{p \mid q \equiv p' \mid q} \quad \frac{p \equiv p'}{(\nu x)p \equiv (\nu x)p'}$$

$$\begin{aligned} !p &\equiv p \mid !p \\ p \mid \text{nil} &\equiv p \\ p \mid q &\equiv q \mid p \\ (\nu x)(\nu y)p &\equiv (\nu y)(\nu x)p \\ (\nu x)\text{nil} &\equiv \text{nil} \\ (\nu x)(p \mid q) &\equiv (\nu x)p \mid q \quad x \text{ not free in } q \end{aligned}$$

Theory of Pi Calculus



Notes

- The Pi calculus does not have the Church-Rosser property

Recall:

$$\text{WireTap} \mid \text{wire}\langle M \rangle \mid \text{ATT} \rightarrow^* \text{NSA}\langle M \rangle \mid \text{fiber}\langle M \rangle$$

But also:

$$\text{WireTap} \mid \text{wire}\langle M \rangle \mid \text{ATT} \rightarrow^* \text{WireTap} \mid \text{fiber}\langle M \rangle$$

- This captures the non-deterministic nature of concurrency

For Pi-calculus there are

- Type systems
- Equivalences and logics
- Expressiveness results, through encodings of numbers, lists, procedures, objects

Pi Calculus Applications



A number of languages are based on Pi calculus in whole or in part.

- e.g., Pict (Pierce and Turner)

Specification and verification.

- mobile phone protocols
- security protocols

The Pi Calculus and Security




The channels of the Pi calculus have nice built-in properties, such as:

- integrity
- confidentiality (with ν)
- exactly-once semantics
- mobility (channels as first-class values)


These properties are useful in high-level descriptions of security protocols

More detailed descriptions are possible in the spi calculus (= pi calculus + cryptography)

A Typical Security Protocol



Establishment and use of a secret channel:



A and B are two clients

S is an authentication server

cAS and cBS are private channels with the server

cAB is a new channel for the clients

The Security Protocol in Pi Calculus



This protocol is described as follows:

$$A(M) = (\nu cAB) \ cAS\langle cAB \rangle . cAB \ \langle M \rangle$$

$$S = ! (cAS(x) . cBS\langle x \rangle \mid cBS(x) . cAS\langle x \rangle)$$

$$B = cBS(x) . x(y) \ \dots$$

$$\text{System}(M) = (\nu cAS) (\nu cBS) \ A(M) \mid S \mid B$$

▸ where ... represents what B does with the message it receives

Some Security Properties



An authenticity property

- For all N , if B receives N then A sent N to B

A secrecy property

- An outsider cannot tell $\text{System}(M)$ apart from $\text{System}(N)$, unless B reveals some part of A 's message

Both of these properties can be formalized and proved in the Pi calculus

The secrecy property can be treated via a simple type system